

TwoPaCo

An efficient algorithm to build the compacted de Bruijn graph
from many complete genomes

Gaspare Ferraro

University of Pisa
Department of Computer Science

Pisa, 29 May 2018



Parte I

Introduction



A pan-genomic algorithm



De Bruijn graph



Compacted de Bruijn graph



Junctions



The problem



Parte II

The algorithm



Naive algorithm

- Store all $(k + 1)$ -mers in a hash table
- For each k -mers query the possible edge
- If only 1 in and 1 out edge, unmark as a junction

Algorithm 1: FILTER-JUNCTIONS

Input : $S = \{s_1, \dots, s_m\}$ genome sequences

k integer, size of k -mers

E empty set data structure

C Candidate set of junctions (**naively all positions are marked**)

Output: A reduce candidate set of junctions C

```

1 foreach  $s \in S$  do
2   for  $1 \leq i < |s| - k$  do
3     if  $C[s, i] = \text{marked}$  then
4        $E \leftarrow E \cup s[i..i+k] \cup s[i-1..i+k-1]$     ▷ Store all  $(k+1)$ -mers

5 foreach  $s \in S$  do
6   for  $1 \leq i < |s| - k$  do
7     if  $C[s, i] = \text{marked}$  then
8        $(in, out) \leftarrow (0, 0)$     ▷ Count in/out edges
9       foreach  $c \in \{A, C, G, T\}$  do
10        if  $v \cdot c \in E$  then
11           $in \leftarrow in + 1$ 
12        if  $c \cdot v \in E$  then
13           $out \leftarrow out + 1$ 
14        if  $(in, out) = (1, 1)$  then
15           $C[s, i] = \text{unmarked}$     ▷ surely not a junction
16 return  $C$ 

```



The memory issue

First part of the naive algorithm:

```

foreach  $s \in S$  do
  for  $1 \leq i < |s| - k$  do
    if  $C[s, i] = \textit{marked}$  then
       $E \leftarrow E \cup s[i..i + k] \cup s[i - 1..i + k - 1]$ 

```

We don't really need and, in almost all practical cases,
we can't store all the possible $(k + 1)$ -mers.

Mainly because **only a little percentual** of them are
junction in the de Bruijn graph.



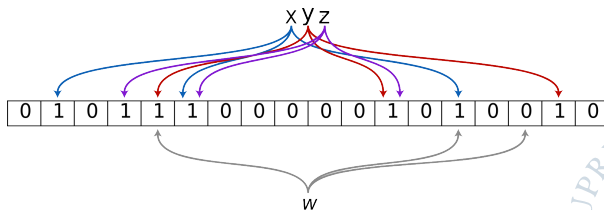
Bloom filter

A space-efficient probabilistic hash table

Bitmap V of size b , h hash functions $f_0, f_1, \dots, f_{h-1} : U \rightarrow [0, b-1]$

insertion(x) $\rightarrow V[f_i(x)] = 1, \forall 0 \leq i < h$

contains(x) \rightarrow **probably yes** if $V[f_i(x)] = 1, \forall 0 \leq i < h$



Probability of false positive, after n insertion: $p_{FP} \simeq (1 - e^{-hn/m})^h$

Two Pass version

- First pass: Select a set of junction candidates by insert all the $(k + 1)$ -mers in a bloom filter of choosing size
- Second pass: Filter out the false positive by storing the reduce sets of $(k + 1)$ -mers in an hash table

Algorithm 2: FILTER-JUNCTIONS-TWO-PASS

Input : strings $S = \{s_1, \dots, s_m\}$ genome sequences

integer k , size of k -mers

integer b , size of bloom filter

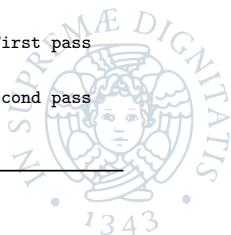
Candidate set of junctions C_{in} (**naively all positions are marked**)

Output: A reduce candidate set of junctions C_{out}

- 1 $F \leftarrow$ empty bloom filter of size b
 - 2 $C_{temp} \leftarrow \text{FILTER-JUNCTIONS}(S, k, F, C_{in})$
 - 3 $H \leftarrow$ empty hash table
 - 4 $C_{out} \leftarrow \text{FILTER-JUNCTIONS}(S, k, H, C_{in})$
 - 5 **return** C_{out}
-

▷ First pass

▷ Second pass



The memory issue²

How much memory do we use now?

- First pass: Bloom filter of size b (of our decision)
- Second pass: Hash table containing $(k + 1)$ -mers of junction candidates

We don't know the possible size of the hash table in the second pass.
What if the hash table is not small enough?

Solution: Split the input k -mers in chunks and analyze them in multiple rounds.



k -mers splitting

Algorithm 3: ROUND-SPLITTING

Input : strings $S = \{s_1, \dots, s_m\}$ genoma sequences

integer k , size of k -mers

integer b , size of bloom filter

integer l , number of rounds

Output: $(V_0, V_1, \dots, V_{l-1})$ chunks of k -mers

```

1  $V_0 \leftarrow \emptyset, V_1 \leftarrow \emptyset, \dots, V_{l-1} \leftarrow \emptyset$ 
2  $c_0 \leftarrow 0, c_1 \leftarrow 0, \dots, c_{q-1} \leftarrow 0$ 
3  $F \leftarrow$  empty Bloom filter of size  $b$ 
4 foreach  $s \in S$  do
5   for  $1 \leq i < |s| - k$  do
6     if  $s[i..i+k-1]$  not in  $F$  then
7       Insert  $s[i..i+k-1]$  in  $F$ 
8        $c_{f(s[i..i+k-1])} \leftarrow c_{f(s[i..i+k-1])} + 1$ 
9  $T \leftarrow \sum_{0 \leq i < q} c_i / l$ 
10  $acc \leftarrow 0, idx \leftarrow 0$ 
11 for  $0 \leq i < q$  do
12    $V_{idx} = V_{idx} \cup \{i\}$ 
13    $acc \leftarrow acc + c_i$ 
14   if  $acc > T$  then
15      $acc \leftarrow 0$ 
16      $idx \leftarrow idx + 1$ 
17 return  $(V_0, V_1, \dots, V_{l-1})$ 

```



Multiple rounds: dealing with memory restrictions

Algorithm 4: TwoPaCo

Input : strings $S = \{s_1, \dots, s_m\}$ genome sequences

integer k , size of k -mers

integer b , size of bloom filter

integer l , number of rounds

Output: C_{final} all the junctions in the compacted de Bruijn graph

```

1  $V_0 \leftarrow \emptyset, V_1 \leftarrow \emptyset, \dots, V_{l-1} \leftarrow \emptyset$ 
2  $c_0 \leftarrow 0, c_1 \leftarrow 0, \dots, c_{q-1} \leftarrow 0$ 
3  $F \leftarrow$  empty Bloom filter of size  $b$ 
4 foreach  $s \in S$  do
5   for  $1 \leq i < |s| - k$  do
6     if  $s[i..i+k-1]$  not in  $F$  then
7       Insert  $s[i..i+k-1]$  in  $F$ 
8        $c_f(s[i..i+k-1]) \leftarrow c_f(s[i..i+k-1]) + 1$ 
9  $T \leftarrow \Sigma_x y$ 
10  $acc \leftarrow 0, idx \leftarrow 0$ 
11 for  $0 \leq i < q$  do
12    $V_{idx} = V_{idx} \cup i$ 
13    $acc \leftarrow acc + c_i$ 
14   if  $acc > T$  then
15      $acc \leftarrow 0$ 
16      $idx \leftarrow idx + 1$ 
17  $C_{init} \leftarrow$  Boolean array with every position unmarked
18 for  $0 \leq i < l$  do
19    $C_i \leftarrow$  mark every position of  $C_{init}$  that starts a  $k$ -mer with hash in  $V_i$ 
20    $C_i \leftarrow \text{FILTER-JUNCTIONS-TWO-PASS}(S, k, b, C_i)$ 
21  $C_{final} = \cup C_i$ 
22 return  $C_{final}$ 

```



Parallelization scheme

As far we focused in reducing memory. What about the time?

TwoPaCo can be easily parallelizable:

- Parallel for
- Concurrent bloom filter
- Concurrent hash table



Parte III

Results



Source code & Dataset

Original source code by medvedev group available on:

<https://github.com/medvedevgroup/TwoPaCo>

Personal implementation available on:

<https://github.com/GaspereG/TwoPaCo>

Dataset for experiments:

- 5 humans (from human reference genome)
- 8 primates (from ...)
- 62 Escherichia coli (from ...)
- 100 simulated human (from ...)



Memory complexity

The memory complexity is the maximum among the first and the second pass of TwoPaCo

- First pass: insert all k -mers in a bloom filter of size b
- Second pass: store all **junction candidates** in a hash table

How many k -mers?

$\mathcal{O}(m)$, where $m = \sum_{s \in S} |s|$ is the total input size

How many junction candidates?

- Real junction, J
- False positive induced from the bloom filter, FP

Result: $\mathcal{O}(\max\{b, (J + FP)k\})$



Time complexity

The memory complexity is the sum between the first and the second pass of TwoPaCo

- First pass: insert all k -mers in a bloom filter of size b using h hash functions
- Second pass: iterate over all **candidate positions** and query the hash table

How many k -mers?

$\mathcal{O}(m)$, where $m = \sum_{s \in S} |s|$ is the total input size

How many candidate positions?

- Real positions, $|G_c|$
- False positive induced from the bloom filter, FP

Result: $\mathcal{O}(mh + (|G_c| + FP)k)$



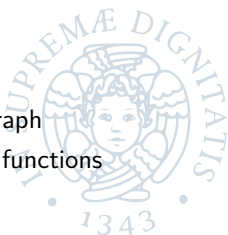
Complexity comparison

State of the art for compressed de Bruijn graph construction:

- Sibelia (Minkin, Patel, Kolmogorov, Vyahhi, Pham, 2013)
- SplitMEM (Marcus, Lee, Schatz, 2014)
- bwt-based (Baier, Beller, Ohlebusch, 2015)

Algorithm	Time complexity	Memory complexity
Sibelia	$\mathcal{O}(m)$	$\mathcal{O}(m)$
SplitMEM	$\mathcal{O}(m \log g)$	$\mathcal{O}(m + G_c)$
bwt-based	$\mathcal{O}(m)$	$\mathcal{O}(m)$
TwoPaCo	$\mathcal{O}(mh + (G_c + FP)k)$	$\mathcal{O}(\max\{b, (J + FP)k\})$

- $m = \sum_{s \in S} |s|$, total input size
- $g = \max_{s \in S} |s|$, size of the biggest genoma
- J and $|G_c|$, number of vertex and edge in the de Bruijn Graph
- b and h , size of the bloom filter table and number of hash functions
- FP , number of false positives in first pass



Running time comparison

What are the practical performances of TwoPaCo against the state of the art?

Dataset	Sibelia	SplitMEM	bwt-based	TwoPaCo	
				1 thread	15 thread
62 E.coli (k=25)	0 (12.2)	70 (178.0)	8 (0.85)	4 (0.16)	2 (0.39)
62 E.coli (k=100)	8 (7.6)	67 (178.0)	8 (0.50)	4 (0.19)	2 (0.39)
7 humans (k=25)	-	-	867 (100.30)	436 (4.40)	63 (4.84)
7 humans (k=100)	-	-	807 (46.02)	317 (8.42)	57 (8.75)
8 primates (k=25)	-	-	-	914 (34.36)	111 (34.36)
8 primates (k=100)	-	-	-	756 (56.06)	101 (61.68)
50 humans (k=25)	-	-	-	-	705 (69.77)
50 humans (k=100)	-	-	-	-	927 (70.21)

Running times in minutes and memory usage in gigabytes in parenthesis

- - on Sibelia = out of time
- - on SplitMEM = out of memory
- - on bwt-based = out of memory
- - on TwoPaCo = experiment not done



Fixed memory

How many rounds do we need to compress the graph without exceeding a given memory threshold?

Threshold	Used memory	Bloom filter size	Running time	Rounds
10GB	8.62GB	8.59GB (2^{33})	4h	1
8GB	6.73GB	4.29GB (2^{32})	7h	3
6GB	5.98GB	4.29GB (2^{32})	9h	4
4GB	3.51GB	2.14GB (2^{31})	11h	6

Experiments on 5 simulated human genomes, $k = 25$, 8 threads.

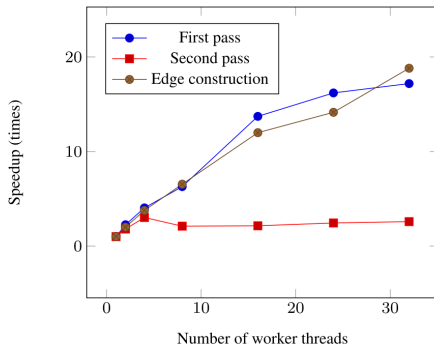
Result: we can trade-off memory for time.



Parallel scalability

How does the performance of TwoPaCo improve to the increasing of working threads?

Parallel scalability



- First pass: great improvement thanks to concurrent bloom filter
- Second pass: slight improvement due to race-conditions on the hash table
- Edge construction: great improvement thanks to k -mers independency

Bloom Filter false positive

Is the Bloom Filter really efficient in reducing junction candidates?

Dataset	k	Junction candidates		
		Initial	First pass	Second pass
62 E.coli	25	310 157 564 (100%)	24 649 489 (7.94%)	24 572 562 (7.92%)
62 E.coli	100	310 157 489 (100%)	22 848 018 (7.36%)	9 492 091 (3.06%)
7 humans	25	21 201 290 922 (100%)	3 489 946 013 (16.46%)	2 974 098 154 (14.02%)
7 humans	100	21 201 290 847 (100%)	1 374 287 870 (6.48%)	188 224 214 (0.88%)
8 primates	25	24 540 556 921 (100%)	5 423 003 377 (22.09%)	5 401 587 503 (22.01%)
8 primates	100	24 540 556 846 (100%)	1 174 160 336 (4.78%)	502 441 107 (2.04%)

- Initial: total number of k -mers in dataset
- First pass: number of junction candidates (using a bloom filter)
- Second pass: real number of junction (using an hash table)



References

[1] Minkin, I., Pham, S., Medvedev, P. (2016).

TwoPaCo: An efficient algorithm to build the compacted de Bruijn graph from many complete genomes

[2] Minkin, I., Patel, A., Kolmogorov, M., Vyahhi, N., Pham, S. (2013).

Sibelia: a scalable and comprehensive synteny block generation tool for closely related microbial genomes.

[3] Marcus, S., Lee, H., Schatz, M. C. (2014).

SplitMEM: a graphical algorithm for pan-genome analysis with suffix skips.

[4] Baier, U., Beller, T., Ohlebusch, E. (2015).

Graphical pan-genome analysis with compressed suffix trees and the **Burrows-Wheeler transform**.



Conclusion

Thanks for your attention!

Questions?

